

Data Transfer Matters of Real-Time GPU Computing

Yusuke Fujii, Takuya Azumi, and Nobuhiko Nishio
College of Information Science and Engineering
Ritsumeikan University

Shinpei Kato and Masato Edahiro
Department of Information Science
Nagoya University


Abstract—Graphics processing units (GPUs) embrace many-core compute devices. In GPU-accelerated systems, massively parallelized compute kernels are offloaded to the GPU. This heterogeneous nature of GPU computing poses non-trivial data transfer issues especially for low-latency real-time systems. A detailed mechanism of data transfers for the GPU must be well studied. In this paper, we investigate possible data transfer methods for the GPU in the state of the art. We also provide open-source implementations of these methods to compare their advantage and disadvantage in performance using real-world systems. Our experimental results show that the hardware-based direct memory access (DMA) and the I/O read and write methods are the most effective, while on-chip microcontrollers inside the GPU are useful to reduce the data transfer times of concurrent multiple data streams. Our findings also include that CPU priorities can protect the performance of GPU data transfers in the context of real-time systems.

Keywords—GPUs; Data Transfer; Latency; Performance


I. INTRODUCTION

Graphics processing units (GPUs) are becoming powerful many-core compute devices. For example, NVIDIA GPUs integrate more than 1,500 processing cores on a single chip and the peak double-precision performance exceeds 1 TFLOPS while sustaining thermal design power (TDP) in the same order of magnitude as traditional multicore CPUs [17]. This rapid growth of GPUs is due to recent advances in the programming model, often referred to as general-purpose computing on GPUs (GPGPU). Data-parallel and compute-intensive applications receive significant performance benefits using GPGPU. Currently a main application of GPGPU is supercomputing [20] but there are more and more emerging applications in different fields. Examples include plasma control [8], autonomous driving [14], software routing [6], encrypted networking [7], and storage management [2], [5], [13], [19]. This broad range of applications raises the need of further developing GPU technology to support scalability of emerging data-parallel and compute-intensive applications.

One of the greatest challenges of GPU computing is to integrate real-time systems. So far the real-time systems community has addressed resource management issues of GPUs [1], [3], [4], [10]–[12]. The main contribution of these work is the scheduling of compute kernels and data transfers associated with the GPU. On the other hand, the basic performance and latency issues of GPUs are not well



(a) Host to Device



(b) Device to Host

Figure 1. Performance of DMA and I/O read/write for the NVIDIA GPU.

explored in the literature. Given that compute kernels are offloaded to the GPU, their performance and latency are more dominated by compiler and hardware technology. However, an optimization of data transfers must be complemented by system software due to constraints of PCIe devices [13]. Data transfers are also interfered by compelling workload on the CPU, while compute kernels are protected within the GPU. These data transfer issues must be well understood and addressed in order to build predictable soft real-time, if not hard real-time, systems augmented with GPUs.

Data transfer issues are particularly important for low-latency GPU computing. Kato *et al.* demonstrated that the data transfer is a dominant property for GPU-accelerated plasma control systems [8]. This is a specific application where data must be transferred between the sensor/actuator devices and the GPU at a high-rate but is a good example presenting the impact of data transfers in GPU computing. Since emerging applications augmented with GPUs may demand a similar performance requirement, a more general study on the GPU data transfer is desired.

Figure 1 illustrates the average data transfer times of hardware-based direct memory access (DMA) and memory-mapped I/O read and write access performed on an NVIDIA GeForce GTX 480 graphics card using an open-source Linux driver [13]. Apparently the performance characteristics are not identical between the host-to-device and the device-to-host directions. A very elementary issue of this performance difference has been discussed [13], but there is no clear conclusion on what methods can optimize the data transfer performance, what different methods are available, and what is their implication for real-time systems. Currently we pray that the black-box data transfer methods provided by vendors' proprietary software are well optimized, because the hardware details of GPUs are not disclosed to the public. However, real-time systems must build on predictable disciplines. A better understanding of the data transfer performance for the GPU is an essential piece of work to facilitate research on GPU-accelerated heterogeneous real-time systems.

Contribution: This paper explores data transfer methods for GPU computing and provides their detailed empirical comparison. We unveil several data transfer methods for the GPU in addition to well-known DMA and I/O read and write approaches. The advantage and disadvantage of these methods are also discussed. The contribution of this paper clarifies how to design and implement optimized data transfer methods especially in the context of real-time constraints. We believe that this is a useful contribution to advance real-time technology with GPUs and emerging heterogeneous compute devices.

Organization: The rest of this paper is organized as follows. Section II presents the assumption and terminology behind this paper. Section III provides an open investigation of data transfer methods for GPU computing. Section IV compares the performances of the investigated data transfer methods in the context of real-time systems. We provide our concluding remarks in Section V.

II. ASSUMPTION AND TERMINOLOGY

We assume the Compute Unified Device Architecture (CUDA) for GPU programming [18]. A unit of code that is individually launched on the GPU is called a *kernel*. The kernel is composed of multiple *threads* that execute the code in parallel. A unit of threads that are co-scheduled by hardware is called a *block*, while a collection of blocks for the corresponding kernel is called a *grid*. The maximum number of threads that can be contained by an individual block is defined by the GPU architecture.

CUDA uses a set of an application programming interface (API) functions to control the GPU. A CUDA program often follows (i) allocate space to the device memory, (ii) copy input data to the allocated device memory space, (iii) launch the program on the GPU, (iv) copy output data back to the host memory, and (v) free the allocated device




Figure 2. Standard DMA.

memory space. The scope of this paper is related to (ii) and (iv). In particular, we use the `cuMemcpyHtoD()` and the `cuMemcpyDtoH()` functions of the CUDA Driver API, which correspond to (ii) and (iv) respectively. Since an open-source implementation of these functions is available with Gdev [13], we modify them to accommodate the data transfer methods investigated in this paper.

In order to focus on the performance of data transfers between the host and the device memory, we allocate a data buffer to the pinned host memory rather than the typical heap allocated by `malloc()`. This pinned host memory space is mapped to the PCIe address and is never swapped out. It is also accessible to the GPU directly.

Our computing platform contains a single set of the CPU and the GPU. Although we restrict our attention to CUDA and the GPU, the notion of the investigated data transfer methods is well applicable to other heterogeneous compute devices. GPUs are currently well-recognized forms of the heterogeneous compute devices, but emerging alternatives include the Intel Many Integrated Core (MIC) and the AMD Fusion technology. The programming models of these different platforms are almost identical in that the CPU controls the compute devices. Our future work includes an integrated investigation of these different platforms.

III. DATA TRANSFER METHODS

In this section, we investigate data transfer methods for the GPU. A standard data transfer method uses hardware DMA engines integrated on the GPU, while direct data read and write accesses to the GPU device memory are allowed through PCIe BARs. We can also use microcontrollers integrated on the GPU to send and receive data across the host and the device memory. Unfortunately, only a limited piece of these schemes has been studied in the literature. Our investigation and open implementations of these schemes provide a better understanding of data transfer mechanisms for the GPU. Note that we restrict our attention to the NVIDIA GPU architecture, but the concepts of data transfer methods introduced in this paper are mostly applicable to PCIe-connected compute devices.

A. Standard DMA

The most typical method for GPU data transfers is to use standard DMA engines integrated on the GPU. There are two types of such DMA engines for synchronous and




Figure 3. Microcontroller-based data transfer.

asynchronous data transfer operations respectively. We focus on the synchronous DMA engines, which always operate in a sequential fashion with compute engines.

Figure 2 shows a concept of this standard DMA method. To perform this DMA, we write *GPU commands* to an on-board DMA engine. Upon a request of GPU commands, the DMA engine transfers a specified data set between the host and the device memory. Once a DMA transfer starts, it is non-preemptive. This method is often the most effective to transfer a large size of data. Details of the GPU DMA mechanism can be found in previous work [12], [13].

B. Microcontroller-based Data Transfer

The GPU provides on-board microcontrollers to control GPU functional units (compute, DMA, power, temperature, encode, decode, etc.). Albeit tiny hardware, these microcontrollers are available for GPU resource management beyond just controlling the functional units. Each microcontroller supports special instructions to transfer data in the data sections to and from the host or the device memory. The data transfer is offloaded to the microcontroller, *i.e.*, DMA, but is controlled by the microcontroller itself. Leveraging this mechanism, we can provide data communications between the host and the device memory.

Figure 3 shows a concept of this microcontroller-based data transfer method. Since there is no data path to directly copy data between the host and the device memory using a microcontroller, each data transfer is forced to take two hops: (i) the host memory and microcontroller and (ii) the device memory and microcontroller. This is non-trivial overhead but the handling of this DMA is very light-weight as compared to the standard DMA method.

The microcontroller is executing firmware loaded by the device driver. We modify this firmware code to employ an interface for the data communications. The firmware adopts an event-driven design: it invokes only when the device driver sends some command from the CPU. The user program hence first needs to communicate with the device driver to issue such a command. Our implementation uses `ioctl` system calls to achieve this user and device driver communication.

A constraint of this microcontroller approach is that the size of each data transfer is limited by 256 bytes. If the data transfer size exceeds 256 bytes, we have to split a transaction into multiple chunks.




Figure 4. Memory-mapped read and write.

C. Memory-mapped Read and Write

The aforementioned two methods are based on DMA functions. DMA is usually high-throughput but it inevitably incurs overhead in the setup. A small size of data transfers may encounter severe latency problems due to this overhead. One of good examples can be found in the plasma control system [8]. If low-latency is required, direct read and write accesses are more appropriate than hardware-based DMA. Since the GPU as a PCIe-connected device provides memory-mapped regions upon the PCI address space, the CPU can directly access the device memory without using bounce buffers on the host memory.

Figure 4 shows a concept of this memory-mapped read and write method. NVIDIA GPUs as well as most other PCIe devices expose base address registers (BARs) to the system, through which the CPU can access specific areas of the device memory. There are several BARs depending on the target device. NVIDIA GPUs typically provide the following BARs:

- BAR0 Memory-mapped I/O (MMIO) registers.
- BAR1 Device memory aperture (windows).
- BAR2 I/O port or complementary space of BAR1.
- BAR3 Same as BAR2.
- BAR5 I/O port.
- BAR6 PCI ROM.

Often the BAR0 is used to access the control registers of the GPU while the BAR1 makes the device memory visible to the CPU. This direct read and write method is pretty simple. We create virtual address space for the BAR1 region and set its leading address to a specific control register. Thus the BAR1 region can be directly accessed by the GPU using the unified memory addressing (UMA) mode, where all memory objects allocated to the host and the device memory can be referenced by the same address space. Once the BAR1 region is mapped, all we have to do is to manage the page table of the GPU to allocate memory objects from this BAR1 region and call the I/O remapping function supported by the OS kernel to remap the corresponding BAR1 region to the user-space buffer.




Figure 5. Memory-window read and write.

D. Memory-window Read and Write

The BAR0 region is often called memory-mapped I/O (MMIO) space. This is the main control space of the GPU, through which all hardware engines are controlled. Its space is sparsely populated with areas representing individual hardware engines, which in turn are sparsely populated with control registers. The list of hardware engines is architecture-dependent. The MMIO space contains a special subarea for indirect device and host memory accesses, separated from the control registers. This plays a role of windows that make the device memory visible to the CPU in a different way than the BAR1 region.

Figure 5 shows a concept of this memory-window read and write method. To set the memory window, we obtain the leading physical address of the corresponding memory object and set it to a specific control register. By doing so, a limited range of the memory object becomes visible to the CPU through a specific BAR0 region. In case of NVIDIA GPUs, the size of this range is 4MB and the specific BAR0 region begins at 0x70000000 in the MMIO address. Once the window is set, we can read and write this BAR0 region to access data on the device memory.

E. Pinned Host Memory

Yet another approach to GPU data transfers is to use the pinned host memory. As aforementioned, NVIDIA GPUs support UMA. This is due to the graphics address remapping table (GART) employed by the GPU, allowing the system to specify physical host memory addresses directly in the GPU page table as far as they are associated with pinned PCI-mapped pages. Although this approach is also effective to an extent for low-latency GPU computing [8], it is outside of the scope of this paper and we focus on how to access the device memory in this paper. Note that we still use the host memory as buffer space; we do not consider a method that makes the GPU read from and write to this host memory directly through UMA.

IV. EMPIRICAL COMPARISON

We now provide a detailed empirical comparison for the advantage and disadvantage of the data transfer methods


presented in Section III. Our experimental setup is composed of an Intel Core i7 2600 processor and an NVIDIA GeForce GTX 480 graphics card. We use the Linux kernel v2.6.42 and Gdev [13] as the underlying OS and GPGPU runtime/driver software respectively. This set of open-source platforms allows our implementations of the investigated data transfer methods.

The test programs are written in CUDA [18] and are compiled using the NVIDIA CUDA Compiler (NVCC) v4.2 [16]. Note that Gdev is compatible with this binary compiler toolkit. We exclude compute kernels and focus on data transfer functions in this empirical comparison. While the test programs uniformly use the same CUDA API functions, we provide different internal implementations according to the target data transfer methods.


Data streams between the host and the device memory are provided by a single GPU context. Performance interference among multiple GPU contexts is outside the scope of this paper. We evaluate the data transfer performances of both real-time and normal tasks. For the scheduling policies of the Linux kernel, we use `SCHED_FIFO` for real-time tasks while `SCHED_OTHER` for normal tasks, where the real-time tasks are always prioritized over the normal tasks. The real-time capability relies on the default performance of the real-time scheduling class supported by the Linux kernel. We believe that this setup is sufficient for our experiments given that we execute at most one real-time task in the system while multiple data streams may be produced by this task. Overall the scheduling performance issues are outside the scope of this paper.

Henceforth we use the following labels to denote the investigated data transfer methods respectively:

- **DMA** denotes the standard DMA method presented in Section III-A.
- **IORW** denotes the memory-mapped read and write method presented in Section III-C.
- **MEMWND** denotes the memory-window read and write method presented in Section III-D.
- **HUB** denotes the microcontroller-based data transfer method presented in Section III-B, particularly using a *hub* microcontroller designed to broadcast among the actual microcontrollers of graphics processing clusters (GPCs), *i.e.*, CUDA core clusters.
- **GPC** denotes the microcontroller-based data transfer method presented in Section III-B, particularly using a single GPC microcontroller.
- **GPC4** denotes the microcontroller-based data transfer method presented in Section III-B, particularly using four different GPC microcontrollers in parallel. Note that the NVIDIA Fermi architecture [15] provides four GPCs and their microcontrollers can perform individually. Therefore we can split the data transfer into four pieces and make the four microcontrollers work in parallel.




(a) Host to Device




(b) Device to Host

Figure 6. Average performance of each data transfer method with a real-time task.



(a) Host to Device



(b) Device to Host

Figure 7. Worst-case performance of each data transfer method with a real-time task.

A. Basic Performance

Figure 6 shows the average performance of each data transfer method when a real-time task runs alone. Even with this most straightforward setup, there are several interest-


ing observations. Performance characteristics of the host-to-device and the device-to-host communications are not identical at all. In particular, the performance of standard DMA exhibits a 10x difference between the two directions of communications. We believe that this is attributed to hardware capabilities that are not well documented to the public. We also find that the performances of different methods are quite diverse but a combination of DMA and IORW can derive the best performance. For small data, IORW is preferred while DMA is the best performer for large data. The other methods are almost always inferior to either of DMA or IORW. The most significant finding is that IORW becomes very slow for the device-to-host direction. This is due to a design specification of the GPU. It is designed so that the GPU can read data fast from the host computer but compromise write access performance. Another interesting observation is that using multiple GPC microcontrollers to parallelize the data transfer is less effective than a single GPC or HUB controller when the data size is small.

Figure 7 shows the worst-case performance in the same setup as the above experiment. It is important to note that we acquire almost the same results as those shown in Figure 6, though there is some degradation in the performance of DMA for the host-to-device direction. These comparisons lead to some conclusion that we may optimize the data transfer performance by switching between DMA and IORW at an appropriate boundary.


We omit the results of normal tasks in this setup, because they are almost equal to those of real-time tasks shown above. However, real-time and normal tasks behave in a very different manner in the presence of competing workload. This will be discussed in the next subsection.

B. Interfered Performance

Figure 8 shows the average performance of each data transfer method when a real-time task encounters extremely high workload on the CPU. All the methods are successfully protected from performance interference due to the real-time scheduler. Note that DMA shows a better performance than the previous experiments despite the presence of competing workload. This is due to the Linux real-time scheduler feature. DMA is performed by GPU commands, which may impose a suspension on the caller task. In the Linux kernel, a real-time task is awakened in a more responsive manner when switched from a normal task than from an idle task. Therefore when the CPU is fully loaded by normal tasks, a real-time task is more responsive. The same is true for the worst-case performance as shown in Figure 9. We learn from these experiments that CPU priorities can protect the performance of data transfer for the GPU. Note that Gdev uses a polling approach to wait for completions of data transfers. An interrupt approach is also worth being investigated.




(a) Host to Device




(b) Device to Host

Figure 8. Average performance of each data transfer method with a real-time task under high CPU load.




(a) Host to Device




(b) Device to Host

Figure 10. Average performance of each data transfer method with a normal data stream under high CPU load.




(a) Host to Device




(b) Device to Host

Figure 9. Worst-case performance of each data transfer methods with a real-time task under high CPU load.



(a) Host to Device



(b) Device to Host

Figure 11. Worst-case performance of each data transfer method with a normal data stream under high CPU load.

For reference, Figure 10 and 11 show the average and the worst-case performance achieved by a normal task when the CPU encounters extremely high workload same as the preceding experiments. Apparently the data transfer times

increase by orders of magnitude as compared to those achieved by a real-time task. DMA shows a bad performance for small data while it can sustain that performance for large data. This is attributed to the fact that once a DMA


command is fired, the data transfer does not have to compete with CPU workload. The other methods are more or less controlled by the CPU, and thus are more affected by CPU workload. This finding provides a design choice for the system implementation.

From now on, we restrict our attention to a real-time task. We next evaluate the performance of each data transfer method under high memory pressure, creating another task that eats up host memory space. In some sense, this is not a meaningful experiment because we use the pinned host memory space to allocate buffers while the memory pressure is supposed to compel the paged host memory space. Having said that, the memory pressure could still impose indirect interference on real-time tasks [9], [21]. It is worth conducting this experiment. As demonstrated in Figure 12 and 13, the impact of memory pressure on the data transfer performance is negligible. This means that all the data transfer methods investigated in this paper require not much paged host memory space. Otherwise they must be interfered by memory workload.


Figure 14 and 15 show the average and the worst-case performance of each data transfer method respectively, when a misbehaving *hackbench* process coexists. *hackbench* is a tool that generates many processes executing I/O system calls with pipes. The results are all similar to the previous ones. Since the Linux real-time scheduler is now well enhanced to protect a real-time task from such misbehaving workload, these results are obvious and trivial in some sense, but we can lead to a conclusion from a series of the above experiments that *the data transfer methods for the GPU can be protected by the traditional real-time scheduler capability*. This is a useful finding to facilitate an integration of real-time systems and GPU computing.

C. Concurrent Performance

So far we have studied the capabilities of the investigated data transfer methods and their causal relation to a real-time task. We now evaluate the performance of concurrent *two* data streams using different combinations of the data transfer methods as shown in Figure 16 and 17. This is a very interesting result. For the host-to-device direction, the best performance is obtained when both the two tasks use IORW. In this case, the two data streams are not overlapped but are processed in sequential due to the use of the same IORW path. Nonetheless it outperforms the other combinations because the performance of IORW is way higher than the other methods as we have observed in a series of the previous experiments. However, the device-to-host direction shows a different performance. Since IORW becomes slow when the CPU reads the device memory as mentioned in Section IV-A, IORW/IOIW is not the best performer any longer. Instead using the microcontroller(s) provides the best performance until 2MB. This is attributed to the fact that the microcontroller-based data transfer method can be




(a) Host to Device




(b) Device to Host

Figure 12. Average performance of each data transfer method with a real-time task under high memory pressure.




(a) Host to Device




(b) Device to Host

Figure 13. Worst-case performance of each data transfer method with a real-time task under high memory pressure.

overlapped with any other data transfer methods. From 16B to 16KB, a combination of the microcontroller and IORW is the fastest, while that of the microcontroller and DMA is the fastest from 16KB to 2MB. Note that from 16B




(a) Host to Device




(b) Device to Host

Figure 14. Average performance of each data transfer method with a real-time task in the presence of hackbench.




(a) Host to Device




(b) Device to Host

Figure 16. Average performance of combinations of the data transfer methods for concurrent real-time tasks.




(a) Host to Device




(b) Device to Host

Figure 15. Worst-case performance of each data transfer method with a real-time task in the presence of hackbench.



(a) Host to Device



(b) Device to Host

Figure 17. Worst-case performance of combinations of the data transfer methods for concurrent real-time tasks.

to 16KB the performance is aligned with the slow IORW curve. Therefore the choice of HUB, GPC, and GPC4 does not really matter to the performance. However, from 16KB to 2MB the performance is improved by using four

microcontrollers in parallel (*i.e.*, GPC4), since DMA is faster than the microcontroller and thereby the performance is aligned with the microcontroller curve. We learn from this experiment that the microcontroller is useful to overlap

concurrent data streams with DMA or IORW, and using multiple microcontrollers in parallel can further improve the performance of concurrent data transfers.

V. CONCLUSION

REFERENCES

- [1] C. Basaran and K-D. Kang. Supporting Preemptive Task Executions and Memory Copies in GPGPUs. In *Proc. of the Euromicro Conference on Real-Time Systems*, pages 287–296, 2012.
- [2] P. Bhatotia, R. Rodrigues, and A. Verma. Shredder: GPU-accelerated incremental storage and computation. In *Proc. of the USENIX Conference on File and Storage Technologies*, 2012.
- [3] G. Elliott and J. Anderson. Globally Scheduled Real-Time Multiprocessor Systems with GPUs. *Real-Time Systems*, 48(1):34–74, 2012.
- [4] G. Elliott and J. Anderson. Robust Real-Time Multiprocessor Interrupt Handling Motivated by GPUs. In *Proc. of the Euromicro Conference on Real-Time Systems*, pages 267–276, 2012.
- [5] A. Gharaibeh, S. Al-Kiswany, S. Gopalakrishnan, and M. Rippeanu. A GPU-accelerated storage system. In *Proc. of the ACM International Symposium on High Performance Distributed Computing*, pages 167–178, 2010.
- [6] S. Hand, K. Jang, K. Park, and S. Moon. PacketShader: a GPU-accelerated software router. In *Proc. of ACM SIGCOMM*, 2010.
- [7] K. Jang, S. Han, S. Han, S. Moon, and K. Park. SSLShader: cheap SSL acceleration with commodity processors. In *Proc. of the USENIX Conference on Networked Systems Design and Implementation*, 2011.
- [8] S. Kato, J. Aumiller, and S. Brandt. Zero-Copy I/O Processing for Low-Latency GPU Computing. In *Proc. of the IEEE/ACM International Conference on Cyber-Physical Systems*, 2013 (to appear).
- [9] S. Kato, Y. Ishikawa, and R. Rajkumar. CPU Scheduling and Memory Management for Interactive Real-Time Applications. *Real-Time Systems*, 2011.
- [10] S. Kato, K. Lakshmanan, Y. Ishikawa, and R. Rajkumar. Resource Sharing in GPU-accelerated Windowing Systems. In *Proc. of the IEEE Real-Time and Embedded Technology and Applications Symposium*, pages 191–200, 2011.
- [11] S. Kato, K. Lakshmanan, A. Kumar, M. Kelkar, Y. Ishikawa, and R. Rajkumar. RGEM: A Responsive GPGPU Execution Model for Runtime Engines. In *Proc. of the IEEE Real-Time Systems Symposium*, pages 57–66, 2011.
- [12] S. Kato, K. Lakshmanan, R. Rajkumar, and Y. Ishikawa. TimeGraph: GPU Scheduling for Real-Time Multi-Tasking Environments. In *Proc. of the USENIX Annual Technical Conference*, 2011.
- [13] S. Kato, M. McThrow, C. Maltzahn, and S. Brandt. Gdev: First-Class GPU Resource Management in the Operating System. In *Proc. of the USENIX Annual Technical Conference*, 2012.
- [14] M. McNaughton, C. Urmson, J. Dolan, and J-W. Lee. Motion Planning for Autonomous Driving with a Conformal Spatiotemporal Lattice. In *Proc. of the IEEE International Conference on Robotics and Automation*, pages 4889–4895, 2011.
- [15] NVIDIA. NVIDIA’s next generation CUDA computer architecture: Fermi. http://www.nvidia.com/content/PDF/fermi_white_papers/NVIDIA_Fermi_Compute_Architecture_Whitepaper.pdf, 2009.
- [16] NVIDIA. CUDA TOOLKIT 4.2. <http://developer.nvidia.com/cuda/cuda-downloads>, 2012.
- [17] NVIDIA. NVIDIA’s next generation CUDA computer architecture: Kepler GK110. <http://www.nvidia.com/content/PDF/kepler/NVIDIA-Kepler-GK110-Architecture-Whitepaper.pdf>, 2012.
- [18] NVIDIA. CUDA Documents. <http://docs.nvidia.com/cuda/>, 2013.
- [19] W. Sun, R. Ricci, and M. Curry. GPUstore: harnessing GPU computing for storage systems in the OS kernel. In *Proc. of Annual International Systems and Storage Conference*, 2012.
- [20] Top500 Supercomputing Sites. <http://www.top500.org/>.
- [21] T. Yang, T. Liu, E.D. Berger, S.F. Kaplan, and J.E-B. Moss. Redline: First Class Support for Interactivity in Commodity Operating Systems. In *Proc. of the USENIX Symposium on Operating Systems Design and Implementation*, pages 73–86, 2008.