

Data Transfer Matters in Low-Latency GPU Computing

Yusuke Fujii, Takuya Azumi, and Nobuhiko Nishio
College of Information Science and Engineering
Ritsumeikan University

Shinpei Kato and Masato Edahiro
Department of Information Science
Nagoya University




Figure 1. Standard DMA.




Figure 2. Microcontroller-based data transfer.

Abstract—

Keywords—Low Latency; Data Transfer; GPUs

I. DATA TRANSFER METHODS

In this section, we investigate data transfer methods for the GPU. A standard data transfer method uses hardware DMA engines integrated on the GPU, while direct data read and write accesses to the GPU device memory are allowed through PCIe BARs. We can also use microcontrollers integrated on the GPU to send and receive data across the host and the device memory. Unfortunately, only a limited piece of these schemes has been studied in the literature. Our investigation and open implementations of these schemes provide a better understanding of data transfer mechanisms for the GPU. Note that we restrict our attention to the NVIDIA GPU architecture, but the concepts of data transfer methods introduced in this paper are mostly applicable to PCIe-connected compute devices.

A. Standard DMA

The most typical method for GPU data transfers is to use standard DMA engines integrated on the GPU. There are two types of such DMA engines for synchronous and asynchronous data transfer operations respectively. We focus on the synchronous DMA engines, which always operate in a sequential fashion with compute engines.

Figure 1 shows a concept of this standard DMA method. To perform this DMA, we write *GPU commands* to an on-board DMA engine. Upon a request of GPU commands, the DMA engine transfers a specified data set between the host and the device memory. Once a DMA transfer starts, it is non-preemptive. This method is often the most effective to transfer a large size of data. Details of the GPU DMA mechanism can be found in previous work [2], [3].

B. Microcontroller-based Data Transfer

The GPU provides on-board microcontrollers to control GPU functional units (compute, DMA, power, temperature, encode, decode, etc.). Albeit tiny hardware, these microcontrollers are available for GPU resource management beyond just controlling the functional units. Each microcontroller supports special instructions to transfer data in the data sections to and from the host or the device memory. The data transfer is offloaded to the microcontroller, *i.e.*, DMA, but is controlled by the microcontroller itself. Leveraging this mechanism, we can provide data communications between the host and the device memory.

Figure 2 shows a concept of this microcontroller-based data transfer method. Since there is no data path to directly copy data between the host and the device memory using a microcontroller, each data transfer needs to take two hops: (i) the host memory and a microcontroller and (ii) the device memory and a microcontroller. This is non-trivial overhead but the handling of this DMA is very light-weight as compared to the standard DMA method.

C. Memory-mapped Read and Write

The aforementioned two methods are based on DMA functions. DMA is usually high-throughput but it inevitably incurs overhead in the setup. A small size of data transfers may encounter severe latency problems due to this overhead. One of good examples can be found in the plasma control system [1]. If low-latency is required, direct read and write accesses are more appropriate than hardware-based DMA. Since the GPU as a PCIe-connected device provides memory-mapped regions upon the PCI address space, the CPU can directly access the device memory without using bounce buffers on the host memory.




Figure 3. Memory-mapped read and write.




Figure 4. Memory-window read and write.

Figure 3 shows a concept of this memory-mapped read and write method. NVIDIA GPUs as well as most other PCIe devices expose base address registers (BARs) to the system, through which the CPU can access specific areas of the device memory. There are several BARs depending on the target device. NVIDIA GPUs typically provide the following BARs:

- BAR0 Memory-mapped I/O (MMIO) registers.
- BAR1 Device memory aperture (windows).
- BAR2 I/O port or complementary space of BAR1.
- BAR3 Same as BAR2.
- BAR5 I/O port.
- BAR6 PCI ROM.

Often the BAR0 is used to access the control registers of the GPU while the BAR1 makes the device memory visible to the CPU. This direct read and write method is pretty simple. We create virtual address space for the BAR1 region and set its leading address to a specific control register. Thus the BAR1 region can be directly accessed by the GPU using the unified memory addressing (UMA) mode, where all memory objects allocated to the host and the device memory can be referenced by the same address space. Once the BAR1 region is mapped, all we have to do is to manage the page table of the GPU to allocate memory objects from this BAR1 region and call the I/O remapping function supported by the OS kernel to remap the corresponding BAR1 region to the user-space buffer.

D. Memory-window Read and Write

The BAR0 region is often called memory-mapped I/O (MMIO) space. This is the main control space of the GPU, through which all hardware engines are controlled. Its space is sparsely populated with areas representing individual hardware engines, which in turn are sparsely populated with control registers. The list of hardware engines is architecture-dependent. The MMIO space contains a special subarea for indirect device and host memory accesses, separated from the control registers. This plays a role of windows that make the device memory visible to the CPU in a different way than the BAR1 region.

Figure 4 shows a concept of this memory-window read and write method. To set the memory window, we obtain the leading physical address of the corresponding memory object and set it to a specific control register. By doing so, a limited range of the memory object becomes visible to the CPU through a specific BAR0 region. In case of NVIDIA GPUs, the size of this range is 1MB and the specific BAR0 region begins at 0x70000000 in the MMIO address. Once the window is set, we can read and write this BAR0 region to access data on the device memory.

E. Pinned Host Memory

Yet another approach to GPU data transfers is to use the pinned host memory. As aforementioned, NVIDIA GPUs support UMA. This is due to the graphics address remapping table (GART) employed by the GPU, allowing the system to specify physical host memory addresses directly in the GPU page table as far as they are associated with pinned PCI-mapped pages. Although this approach is also effective to an extent for low-latency GPU computing [1], it is not within the scope of this paper and we focus on how to access the device memory in this paper.

II. EMPIRICAL COMPARISON

We now provide a detailed empirical comparison for the advantage and disadvantage of the data transfer methods presented in Section I. Our experimental setup is composed of an Intel Core i7 2600 processor and an NVIDIA GeForce GTX 480 graphics card. We use the Linux kernel v2.6.42 and Gdev [3] as the underlying OS and GPGPU runtime/driver software respectively. This set of open-source platforms allows our implementations of the investigated data transfer methods.

The test programs are written in CUDA [6] and are compiled using the NVIDIA CUDA Compiler (NVCC) v4.2 [5]. Note that Gdev is compatible with this binary compiler toolkit. We exclude compute kernels and focus on data transfer functions in this empirical comparison. While the test programs uniformly use the same CUDA API functions, we provide different internal implementations according to the target data transfer methods.

Data streams between the host and the device memory are provided by real-time tasks in Linux. These real-time tasks are prioritized over other background tasks running in Linux. The real-time capability relies on the default performance of the real-time scheduling class supported by the Linux kernel. We believe that this setup is sufficient for our experiments given that we execute at most one real-time task in the system while multiple data streams may be produced by this task. Scheduling performance issues are not within the scope of this paper.

Henchforth we use the following labels to denote the investigated data transfer methods respectively:


- **DMA** denotes the standard DMA method presented in Section I-A.
- **IORW** denotes the memory-mapped read and write method presented in Section I-C.
- **MEMWND** denotes the memory-window read and write method presented in Section I-D.
- **HUB** denotes the microcontroller-based data transfer method presented in Section I-B, particularly using a *hub* microcontroller designed to broadcast among the actual microcontrollers of graphics processing clusters (GPCs), *i.e.*, CUDA core clusters.
- **GPC** denotes the microcontroller-based data transfer method presented in Section I-B, particularly using a single GPC microcontroller.
- **GPC4** denotes the microcontroller-based data transfer method presented in Section I-B, particularly using four different GPC microcontrollers in parallel. Note that the NVIDIA Fermi architecture [4] provides four GPCs and their microcontrollers can perform individually. Therefore we can split the data transfer into four pieces and make the four microcontrollers work in parallel.

A. Basic Performance


Figure 5 shows the average performance of each data transfer method when a single data stream performs alone. Even with this most straightforward setup, there are several interesting observations. The performance characterizations of the host-to-device and the device-to-host communications are not identical at all.

REFERENCES

- [1] S. Kato, J. Aumiller, and S. Brandt. Zero-Copy I/O Processing for Low-Latency GPU Computing. In *Proc. of the IEEE/ACM International Conference on Cyber-Physical Systems*, 2013 (to appear).
- [2] S. Kato, K. Lakshmanan, R. Rajkumar, and Y. Ishikawa. TimeGraph: GPU Scheduling for Real-Time Multi-Tasking Environments. In *Proc. of the USENIX Annual Technical Conference*, 2011.
- [3] S. Kato, M. McThrow, C. Maltzahn, and S. Brandt. Gdev: First-Class GPU Resource Management in the Operating System. In *Proc. of the USENIX Annual Technical Conference*, 2012.




(a) Host to Device




(b) Device to Host

Figure 5. Average data transfer times of real-time single streams.




(a) Host to Device




(b) Device to Host

Figure 6. Worst-case data transfer times of real-time single streams.

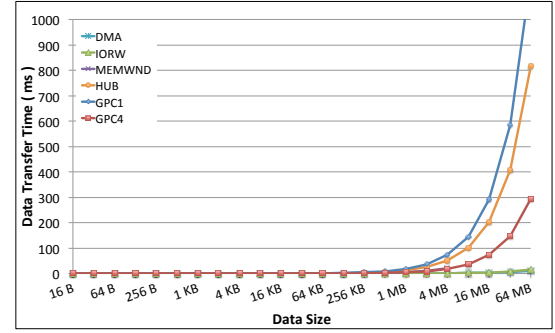


(a) Host to Device




(b) Device to Host

Figure 7. Average data transfer times of real-time single streams under high CPU load.




(a) Host to Device




(b) Device to Host

Figure 8. Worst-case data transfer times of real-time single streams under high CPU load.

- [4] NVIDIA. NVIDIA's next generation CUDA computer architecture: Fermi. http://www.nvidia.com/content/PDF/fermi_white_papers/NVIDIA_Fermi_Compute_Architecture_Whitepaper.pdf, 2009.
- [5] NVIDIA. CUDA TOOLKIT 4.2. <http://developer.nvidia.com/cuda/cuda-downloads>, 2012.
- [6] NVIDIA. CUDA Documents. <http://docs.nvidia.com/cuda/>, 2013.




(a) Host to Device




(b) Device to Host

Figure 9. Average data transfer times of real-time single streams under high memory access load.




(a) Host to Device




(b) Device to Host

Figure 10. Worst-case data transfer times of real-time single streams under high memory access load.




(a) Host to Device




(b) Device to Host

Figure 12. Worst-case data transfer times of real-time single streams in the presence of hackbench.

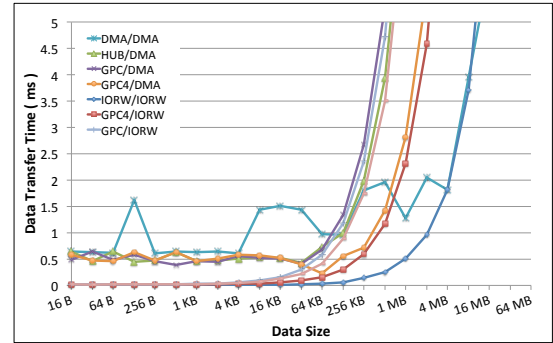


(a) Host to Device




(b) Device to Host

Figure 11. Average data transfer times of real-time single streams in the presence of hackbench.




(a) Host to Device




(b) Device to Host

Figure 13. Average data transfer times of concurrent two streams.

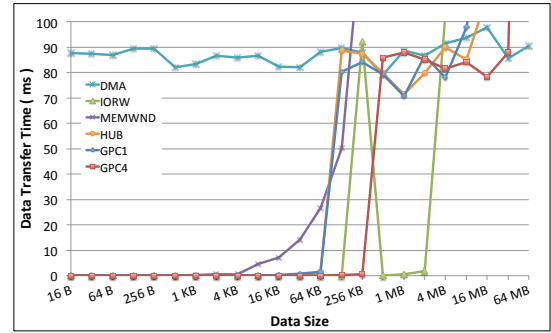


(a) Host to Device




(b) Device to Host

Figure 14. Average data transfer times of single streams.




(a) Host to Device




(b) Device to Host

Figure 16. Average data transfer times of single streams under high CPU load.




(a) Host to Device

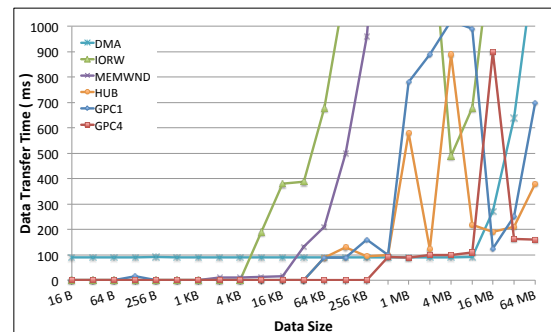


(b) Device to Host

Figure 15. Worst-case data transfer times of single streams.




(a) Host to Device




(b) Device to Host


Figure 17. Worst-case data transfer times of single streams under high CPU load.




(a) Host to Device



(a) Host to Device




(b) Device to Host




(b) Device to Host

Figure 18. Average data transfer times of single streams under high memory access load.


Figure 20. Average data transfer times of single streams in the presence of hackbench.




(a) Host to Device



(a) Host to Device



(b) Device to Host



(b) Device to Host

Figure 19. Worst-case data transfer times of single streams under high memory access load.

Figure 21. Worst-case data transfer times of single streams in the presence of hackbench.